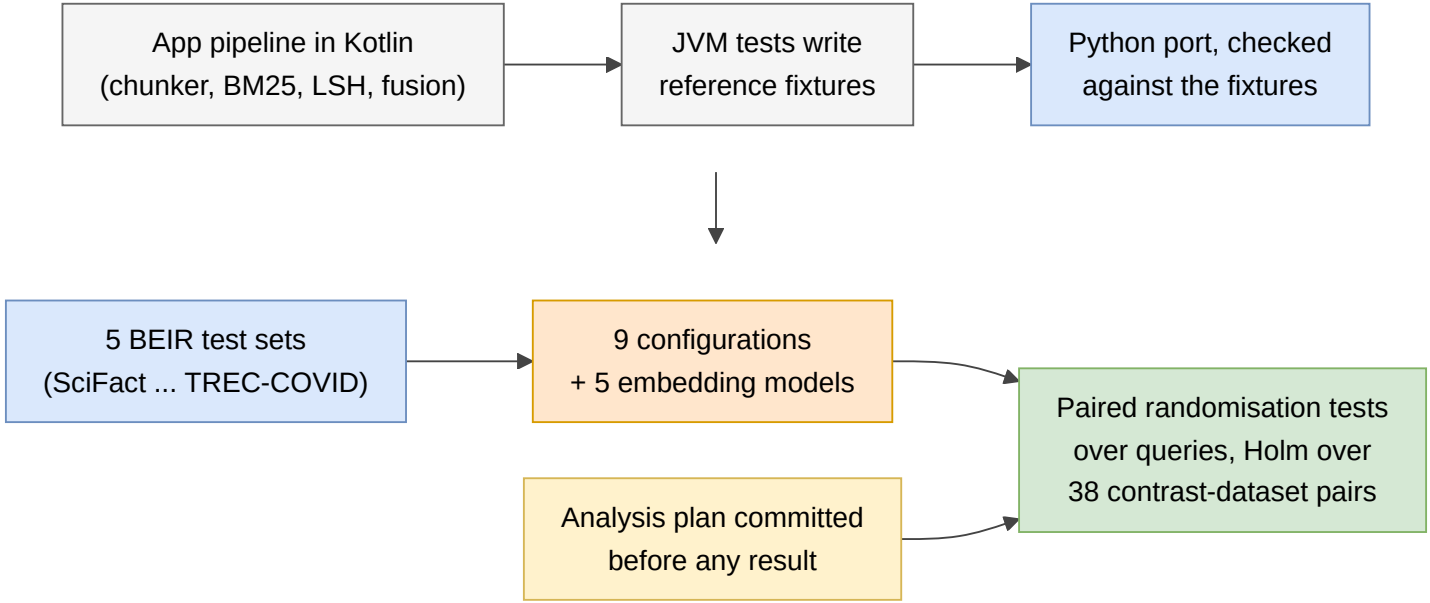
App pipeline in Kotlin
(chunker, BM25, LSH, fusion)
JVM tests write
reference fixtures
Python port, checked
against the fixtures
5 BEIR test sets
(SciFact ... TREC-COVID)
9 configurations
+ 5 embedding models
Analysis plan committed
before any result
Paired randomisation tests
over queries, Holm over
38 contrast-dataset pairs