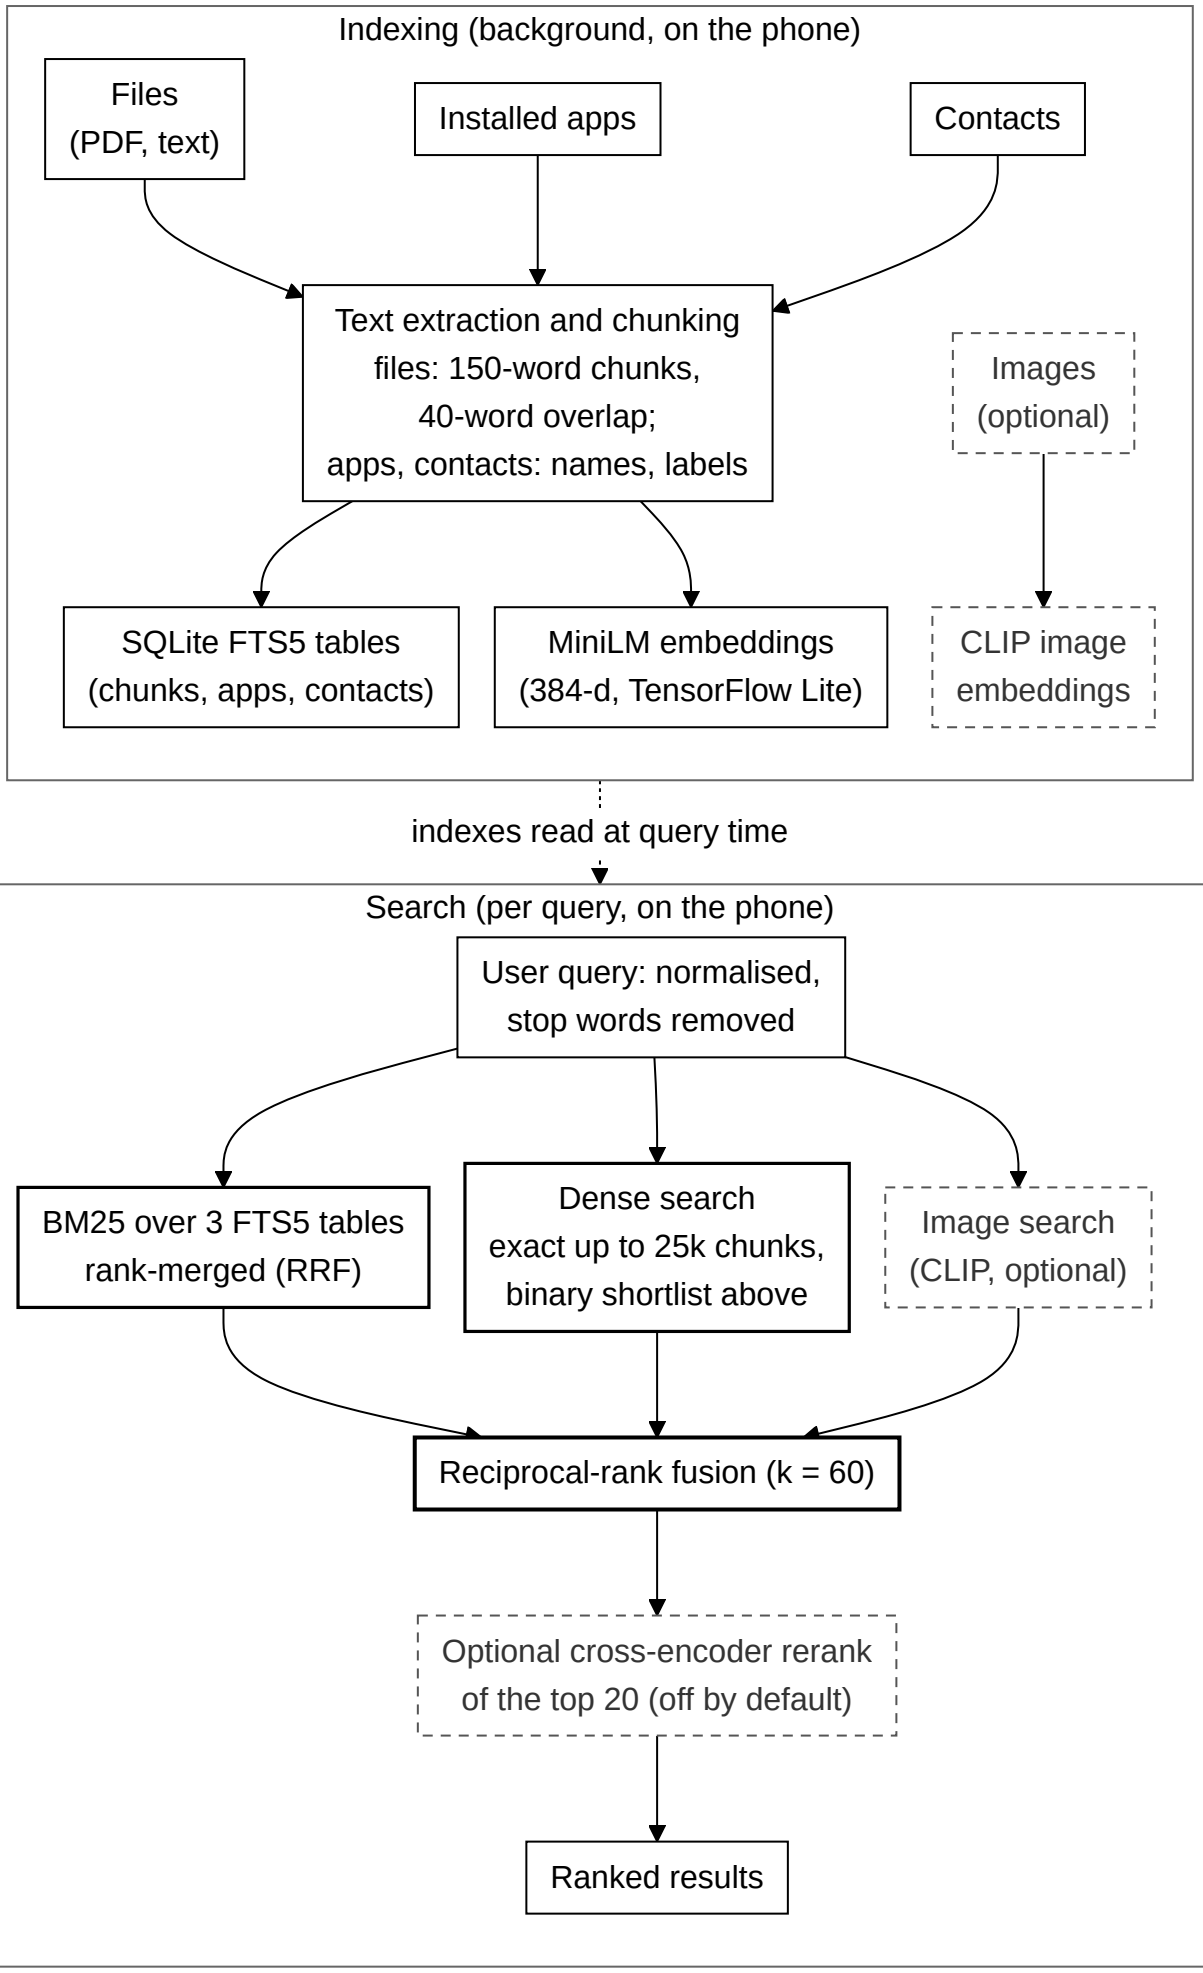
Indexing (background, on the phone)
Files (PDF, text)
Installed apps
Contacts
Text extraction and chunking
files: 150-word chunks,
40-word overlap;
apps, contacts: names, labels
Images (optional)
SQLite FTS5 tables (chunks, apps, contacts)
MiniLM embeddings (384-d, TensorFlow Lite)
CLIP image embeddings
indexes read at query time
Search (per query, on the phone)
User query: normalised, stop words removed
BM25 over 3 FTS5 tables rank-merged (RRF)
Dense search
exact up to 25k chunks,
binary shortlist above
Image search (CLIP, optional)
Reciprocal-rank fusion (k = 60)
Optional cross-encoder rerank of the top 20 (off by default)
Ranked results